

# Comparison of JVM and Dalvik VM

Dainius Jocas

January 30, 2013

## Contents

<b>1</b>	<b>Definition</b>	<b>4</b>
1.1	Definition of JVM as part of Java technology . . . . .	5
1.2	Definition of DVM . . . . .	5
1.3	Comparison of definitions of JVM and DVM . . . . .	5
<b>2</b>	<b>Motivation (Purpose)</b>	<b>5</b>
2.1	Motivation for JVM . . . . .	5
2.2	Motivation for DVM . . . . .	6
2.2.1	J2ME problems . . . . .	7
2.2.2	Platform constraints . . . . .	7
2.3	Comparison of key JVM and DVM motivations . . . . .	7
<b>3</b>	<b>Architecture of virtual machines</b>	<b>8</b>
3.1	What should virtual machine implement? . . . . .	8
3.2	System and process VM's . . . . .	8
3.3	Stack-based VM's . . . . .	8
3.4	Registry-based VM's . . . . .	9
3.5	Main advantages of register-based over stack-based VM's . . . . .	10
<b>4</b>	<b>Executable format</b>	<b>10</b>
4.1	.class file structure . . . . .	10
4.2	.dex file structure . . . . .	11
4.3	Relation between .dex and .class file types . . . . .	12
4.4	Bytecode differences . . . . .	13
<b>5</b>	<b>Code support</b>	<b>13</b>
5.1	Standard Java API libraries that are not supported in Android . . . . .	13
5.2	Android specific libraries . . . . .	15

<b>6</b>	<b>Just-In-Time Compilation</b>	<b>15</b>
6.1	JVM and JIT . . . . .	15
6.2	DVM and JIT . . . . .	16
<b>7</b>	<b>Conclusion</b>	<b>16</b>
<b>A</b>	<b>Comparison of JVM and DVM</b>	<b>18</b>

*Dalvik isn't exactly a household word (at least in my country), and most people wouldn't know a virtual machine if it hit them in the face, but when you tell them you were able to make their existing device work better – run faster, use less battery – they will actually take notice! (Dan Bornstein)[3]*

## Introduction

The popularity of mobile devices is increasing very rapidly. It is considered that the most popular mobile platform for smartphones nowadays is Android[4]. Here we can ask questions such as, why Android is so popular? Does Android meets the requirements for mobile platforms in the best way? To answer those questions we need to take a good look at the requirements for mobile computing and examine the Android approach to meet those requirement for mobile platform.

## Requirements for mobile platform

Smartphone devices comparing to desktop or laptop cousins has various differences. Smartphones have to meet a whole bunch of restrictions that are not very important for desktops. The most important ones are:

- battery power
- passive cooling
- small form factor.

If we take a look at a typical smartphone characteristics: ARM CPU that runs at 500 MHz clock speed, with 128 MB of RAM, no swap space, limited external flash storage (usually NAND), and 1500 mAh battery capacity[11], we can clearly see that we cannot simply use, for example, the same virtual machine as desktop computer with much greater capabilities, because the performance would be unacceptable and we would drain battery in matter of minutes.

## Android approach to meet requirement for mobile platform

Android is an open-source software stack for mobile phones and other devices[6]. Android approach to meet requirements for mobile platform is to provide the whole software stack (e.g. Linux kernel, Java programming language, etc.) which is optimized exactly for mobile platform.

We can see that Android core developers are using technologies that already exist and are very popular. By using popular technologies Android platform can get the attention of lot of skilled developers which is probably the most important ingredient to make a successful mobile platform.

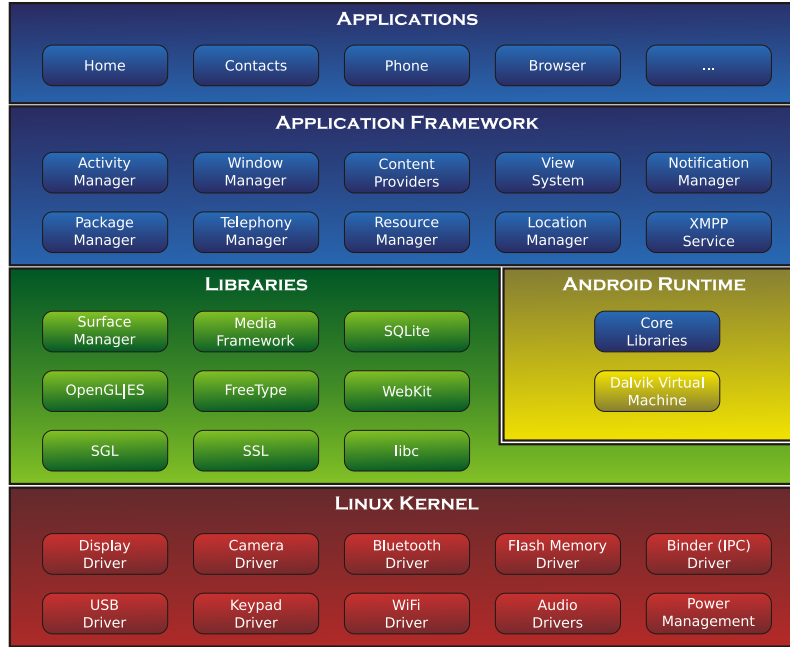


Figure 1: Android system architecture

## The cornerstone for Android

At the core (in terms of importance) of the Android platform we find a Dalvik virtual machine (DVM). Dalvik is the managed runtime used by applications and some system services on Android. Dalvik was originally created specifically for the Android project[5]. DVM promises to run fast and have a small memory footprint on systems that has limited amount of memory and relatively slow CPU. This is exactly what is needed in the modern mobile platform.

DVM is an descendant of Java virtual machine (JVM). The rest of this report will discuss the main differences between the two virtual machines.

## 1 Definition

A virtual machine (VM) is a software implementation of a machine (i.e. a computer) that executes programs like a physical machine[12]. Basic VM parts are: a set of registers, a stack (optional), an execution environment, a garbage-collected heap, a constant pool, a method storage area, an instruction set.

In this section we explore definitions of the JVM and DVM and differences between the two.

## 1.1 Definition of JVM as part of Java technology

The name "Java" for programming language is said to originate from Java coffee, that was consumed in large quantities by programming language authors[16]. Java coffee refers to coffee beans produced in the Indonesian island of Java[15].

Java technology is both Java programming language and Java platform. Java platform has two components: Java virtual machine and The Java Application Programming Interface (API)[7]. At the heart of Java Technology lies the Java virtual machine (JVM) – the abstract computer on which all Java programs run. A Java virtual machine's main job is to load class files and execute the bytecodes they contain. Usually Java programs are written in Java programming language. Also, there are JVM compatible languages, e.g. Scala and Clojure[18].

## 1.2 Definition of DVM

Dalvik virtual machine name originates from fishing village Dalvik in northern part of Iceland[2]. In Dalvik village lived some of ancestors of the author of DVM Dan Bornstein.

DVM is part of Android technology stack which is responsible for running applications. DVM is implemented as a virtual machine that runs programs that are specifically written for Android OS. Usually programs for Android OS are written in Java programming language.

## 1.3 Comparison of definitions of JVM and DVM

From the definition point of view we can conclude that Oracle Inc. had a reason to sue Google Inc. for patent issues because both JVM and DVM are designed to do the same thing - to be a virtual machine that executes programs written in Java programming language. But as it is usual, "devil is in the details" and we'll discover those differences between DVM and JVM in great detail.

# 2 Motivation (Purpose)

In this section we investigate the core motivations to create both JVM and DVM. Also, we take a closer look at technological and business motivations.

## 2.1 Motivation for JVM

The development history of Java platform including JVM begins at Sun Microsystems in December 1990. Primarily motivation was to create an alternative for C/C++ programming languages. In 1994 Java platform was retargeted for World Wide Web. Afterwards the Java platform was developed to address the problems of building software for networked consumer devices. It was designed to support multiple host architectures and to

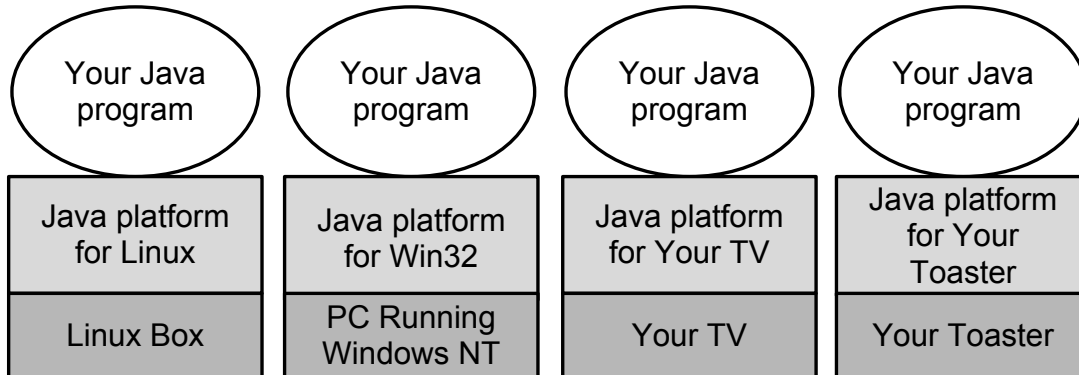


Figure 2: Java programs run on top of Java platform[17].

allow secure delivery of software components. To meet the requirements, compiled code had to survive transport across networks, operate on any client, and assure the client it was safe to run[10].

To meet these requirements one needs tools that empowers developers to make portable programs. In other words, tool that meets the principle of **WORA** (Write Once Run Anywhere). The way to fulfill this principle is to use concept of virtual machines because it gives you an abstraction of an underlying platform. For example, if there is a need to support Java applications on one more hardware and OS platform, the only thing you need to do is to port JVM to that platform and existing applications will work on it. This is possible because JVM executes java byte code which stays hardware- and OS-independent.

## 2.2 Motivation for DVM

As stated in Android homepage:

Our primary purpose is to build an excellent software platform for everyday users[6].

To build such software platform one also needs a virtual machine because there are huge variety of devices that those “everyday users” have. There were many reasons that pushed Google to develop a custom clean-room virtual machine of their own rather than using one of the existing Java Virtual Machines and below we’ll emphasize some of the most important.

### 2.2.1 J2ME problems

While Java technology adoption for desktops is more or less homogenous, for mobile platforms situation is not that attractive. In the days, when Android was just an idea, J2SE (Java 2 Standard Edition) for mobile devices was too “heavy”. Lighter Java version – J2ME (Java 2 Mobile Edition) – was lacking in key features that are very important for modern smartphones such as Bluetooth support or 3D graphics were available only on subset of devices. Because of the presence of “Java Community Process” organization[9], to standardize J2ME would have been long and bureaucratic process. Also, a large part of the reason Java has had limited success in the past on mobile platforms is due to the speed of mobile Java platforms which are not designed to take advantage of today’s smartphone devices.

### 2.2.2 Platform constraints

Android is designed mainly for smartphones and tablets which are constrained by limited resources. Virtual machine that runs in those devices should address those constraints. DVM is designed to run in the device which has:

- Slow CPU: 250-500 MHz;
- Relatively little amount of RAM: total 64 MB, from which 20 MB are available for applications;
- OS without swap space;
- Device is powered by a battery;

When compared to the desktops, these resources are very limited. So, there was a need to rethink the whole software stack and, at the same time, to find a VM solution for mobile platforms.

## 2.3 Comparison of key JVM and DVM motivations

Common motivation for both JVM and DVM is that they are trying to address the problem of running software in heterogeneous platforms which are used by “everyday users”.

Security is always important issue for information systems. It is widely known that VM technology provides isolation between multiple systems running concurrently on the same hardware platform[14].

The difference is that in DVM case, the platform is very resource constrained, when compared to the platform JVM is aiming to.

### 3 Architecture of virtual machines

The term architecture is used here to describe the attributes of the system as seen by the programmer, i.e., **the conceptual structure and functional behavior**, as distinct from the organization of the data flow and controls, the logic design, and the physical implementation[1].

In this section we discuss what VM should implement. Furthermore, we take a brief look at architectural difference of system and process VM's. Also, we describe stack-based and registry-based VM's, the main differences between them, and the consequences of those differences.

#### 3.1 What should virtual machine implement?

A virtual machine should emulate the operations carried out by a physical CPU and thus should ideally encompass the following concepts[3]:

- Compilation of source language into VM specific byte code;
- Data structures to contain instructions and operands;
- A call stack for function call operands;
- An "Instruction Pointer" (IP) pointing to the next instruction to execute;
- A virtual CPU - the instruction dispatcher that:
  - Fetches the next instruction (addressed by the instruction pointer);
  - Decodes the operands
  - Executes the instructions;

#### 3.2 System and process VM's

A system virtual machine provides a complete system platform which supports the execution of a complete operating system (OS). A process virtual machine (also, language virtual machine) is designed to run a single program, which means that it supports a single process[19].

Conceptual difference between system and process VM's is shown in figure 3. The main advantage of using process VM is that if a VM crashes the other processes are not affected.

Both JVM and DVM are process VM's. From this perspective JVM and DVM are similar.

#### 3.3 Stack-based VM's

A stack-based VM implements previously described concepts using the data structure of the stack for memory organization. In this data structure operands are stored. Operands



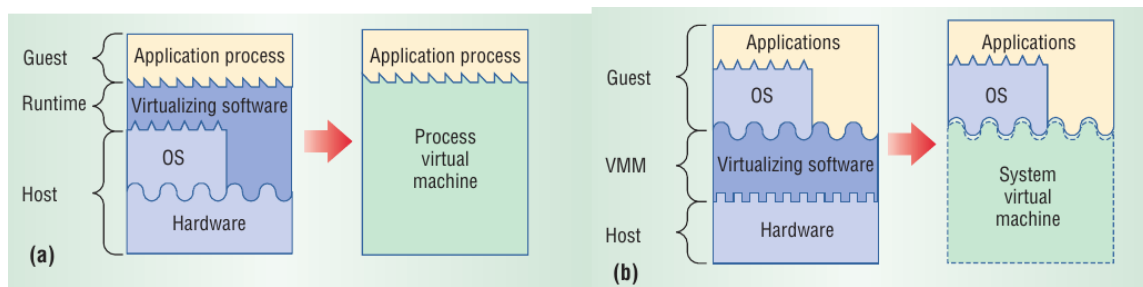


Figure 3: Conceptual architecture of system and process virtual machines[14]. a)

are carried out by popping (POP) data from the stack, processing the popping out (POP) and pushing in (PUSH) back the result in LIFO (Last In First Out) fashion.

The advantage of the stack based model is that the operands are addressed implicitly by the stack pointer (SP). This means that the VM does not need to know the operand addresses explicitly, as calling the stack pointer will give (POP) the next operand. In stack-based VM's, all the arithmetic and logic operations are carried via pushing and popping the operands and results in the stack.

In stack-based VM's adding two numbers usually is carried out in the following manner:

1. POP 20
2. POP 7
3. ADD 20, 7, result
4. PUSH result

Illustration of addition process is in the figure 4.




Figure 4: Addition in stack-based VM architecture.




Figure 5: Addition in register-based VM architecture.

### 3.4 Registry-based VM's

In the register-based VM implementations, the data structure where operands are stored is based on the registers of the CPU. There is no PUSH or POP operations, therefore

the instructions need to contain the addresses (the registers) of the operands. That is the operands for the instructions are explicitly addressed in the instruction, unlike the stack based model where we had a stack pointer (SP) to point to the operand.

Compared to stack-based VM addition is just one operation: **ADD R1, R2, R3**;. Graphical interpretation of the register-based addition process is in the figure 5.

### 3.5 Main advantages of register-based over stack-based VM's

Main advantages of registry-based VMs compared to stack-based VMs that are important for mobile environments:

- a register-based VM reduces the number of instructions to be executed, because one instruction is more expressive;
- the resulting executable code is larger, yet it requires less time to run;
- avoids instruction dispatch, because in general there are fewer instructions;
- avoids unnecessary memory access;
- consume instruction stream efficiently;

Generally speaking, the code size of register-based VM instructions are larger than that of the corresponding stack-based VM instructions. At the same time the instruction number is lower. It means that less CPU usage and then power saving.

Having all operands in the instruction has its benefits; the execution is faster compared to the stack VM, which needs a small calculation to find out the operand, while register VMs just read the registers.

In register VMs, temporary values usually remain in registers. Stack VMs are unable to use a few optimisations too. For example, in the case of common sub-expressions, which are recalculated each time they appear in the code, a register VM can calculate an expression once, and keep that in a register for all future references.

Because code efficiency and power management is of a paramount importance for mobile environment, DVM was implemented as a registry-based VM.

## 4 Executable format

In this section we describe internal file structure of .class and .dex file formats. Also, we discover relation between the two file types. Moreover, we provide analysis of bytecode.

### 4.1 .class file structure

In standard Java environments, Java source code is compiled into Java Bytecode, which is stored within .class files. The .class files are read by the JVM at runtime. Each class in your Java code will result in one .class file. This means that if you have, say, one .java source

file that contains one public class, one static inner class, and three anonymous classes, the compilation process (javac) will output 5 .class files.

An internal structure of .class files[8]:

- A header containing a “magic number” and a version number;
- A constant pool;
- Access rights of the class encoded by a bit mask;
- A list of interfaces implemented by the class;
- A list containing fields of the class;
- A list of methods of the class;
- A list of the class attributes (e.g. exceptions);



Figure 6: .class file logical structure.

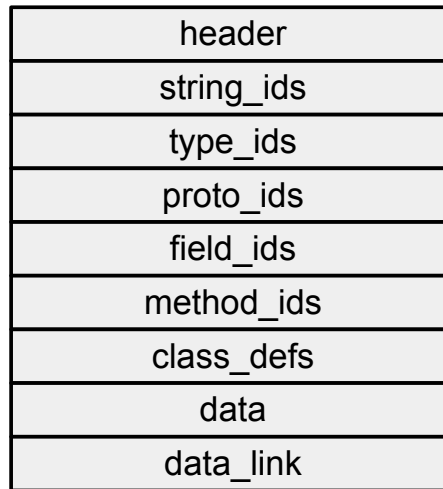


Figure 7: .dex file logical structure.

Typically .class files are distributed in collections. Those collections are .jar files. .jar files are nothing more than zipped collection of .class files.

## 4.2 .dex file structure

Bytecode for DVM is stored in .dex (Dalvik EXecutable) files. Bytecode for .dex files are generated from byte code stored .class files using the **dx** tool. The format of the .dex file is the following [13]:

- Header with “magic number” and version number;

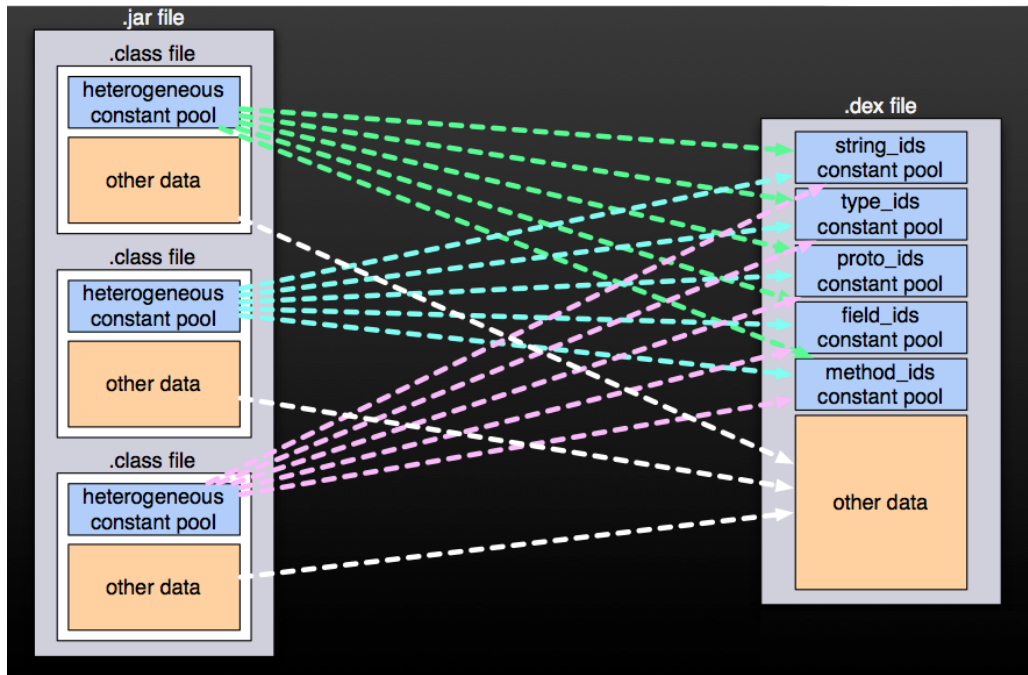


Figure 8: .class files compiled into .dex file.

- A constant pool for each family of elements (strings, fields, methods...);
- Definition of classes;
- Definitions of methods;
- Reserved space for general data;

Dex file has a limitation of 64K method references. Usually, one Android application has one .dex file, which has all the necessary code. But sometimes, big applications can contain more than 64K method references. To get around this limitation, developers can partition part of the program into multiple secondary dex files, and load them at runtime.

### 4.3 Relation between .dex and .class file types

As stated before .dex files are compiled from .class files. This compilation takes all the needed .class files, analyzes them, optimizes the overall structure and puts everything into a .dex file. Basically, .dex file is an optimized collection of .class files, e.g. figure 8. These optimizations leads to a measurable savings in memory consumption. These saving are achieved due to the fact that .dex files avoid repetition of data, and implicit typing and implicit labeling are used.

Some of the optimizations done while constructing .dex files:

- All separate constant pools are consolidated into one shared constant pool. Therefore, minimal repetition, per-type pools (implicit typing), implicit labeling.
- For instance field get/put, replace the field index with a byte offset. Also, merge the boolean / byte / char / short variants into a single 32-bit form (less code in the interpreter means more room in the CPU I-cache).
- Symbolic references from one class file to another are changed to static pointers.
- Inlining special native methods, e.g. `String.length()`.
- Byte-swapping and padding.
- Pruning empty methods.
- Adding auxiliary data, e.g. a hash table for lookups on class name.

#### 4.4 Bytecode differences

The most significant differences between JVM and DVM bytecode:

- Code unit in JVM is one byte, Dalvik - 2 bytes.
- DVM has 230 opcodes while JVM has 225 different opcodes.
- In JVM, the opcodes for integers and single float constants are different, along with long integers and double float constants. While Dalvik implements the same opcode for both, integers and float constants.
- Dalvik bytecode does not have a specific null type. Instead, Dalvik uses a 0 value constant. So, the ambiguous implication of constant 0 should be distinguished properly.
- JVM bytecode uses different opcodes for object reference comparison and null type comparison, while Dalvik simplifies them into one opcode. Thus, the type info of the comparison object must be recovered during decompilation.

## 5 Code support

Usually applications for Android OS are written in Java programming language. Android API supports a relatively large subset of Java Standard Edition 5.0 library but some libraries are left out. Some things were left out because they simply didn't make sense (like printing), and others because better APIs are available that are specific to Android (like user interfaces).

In this section we take a look at which libraries are left out of Standard Java API. Also, we briefly describe what are specific Java libraries for android.

### 5.1 Standard Java API libraries that are not supported in Android

Table 1: Libraries from Java API that are not supported in Android API

Library	Comment
java.applet	Provides the classes necessary to create an applet and the classes an applet uses to communicate with its applet context.
java.awt	Contains all of the classes for creating user interfaces and for painting graphics and images.
java.beans	Contains classes related to developing beans – components based on the JavaBeans™ architecture.
java.lang.management	Provides the management interface for monitoring and management of the Java virtual machine as well as the operating system on which the Java virtual machine is running.
java.rmi	Java Remote Method Invocation (Java RMI) enables the programmer to create distributed Java technology-based to Java technology-based applications, in which the methods of remote Java objects can be invoked from other Java virtual machines*, possibly on different hosts.
javax.accessibility	Defines a contract between user-interface components and an assistive technology that provides access to those components.
javax.activity	Contains Activity service related exceptions thrown by the ORB machinery during unmarshalling.
javax.imageio	The main package of the Java Image I/O API.
javax.management	Provides the core classes for the Java Management Extensions.
javax.naming	Provides the classes and interfaces for accessing naming services.
javax.print	Provides the principal classes and interfaces for the Java™ Print Service API.
javax.rmi	Contains user APIs for RMI-IIOP.
javax.security.auth.kerberos	This package contains utility classes related to the Kerberos network authentication protocol.
javax.security.auth.spi	This package provides the interface to be used for implementing pluggable authentication modules.
javax.security.sasl	Contains class and interfaces for supporting SASL.

Table continues on the next page...

Table 1 – Continued

Library	Comment
javax.swing	Provides a set of "lightweight" (all-Java language) components that, to the maximum degree possible, work the same on all platforms.
javax.transaction	Provides the API that defines the contract between the transaction manager and the various parties involved in a distributed transaction namely : resource manager, application, and application server.
javax.xml (except javax.xml.parsers)	This package contains the core JAX-WS APIs.
org.ietf.*	This package presents a framework that allows application developers to make use of security services like authentication, data integrity and data confidentiality from a variety of underlying security mechanisms like Kerberos, using a unified API.
org.omg.*	Provides the mapping of the OMG CORBA APIs to the Java™ programming language, including the class ORB, which is implemented so that a programmer can use it as a fully-functional Object Request Broker (ORB).
org.w3c.dom.* (sub-packages)	Provides the interfaces for the Document Object Model (DOM) which is a component API of the Java API for XML Processing.

## 5.2 Android specific libraries

In this section we'll take a look which standard Java libraries are unavailable to use in DVM and which libraries are specific to DVM are unavailable to use in JVM.

## 6 Just-In-Time Compilation

Just-in-time compilation (JIT), also known as dynamic translation, is a method to improve the runtime performance of computer programs based on bytecode (virtual machine code). It works translating the original bytecode to the native code, e.g. Java bytecode to assembly code.

In this section we describe JVM relationship with JIT and DVM approach to JIT.

### 6.1 JVM and JIT

On a Java virtual machine implemented in software, the simplest kind of execution engine just interprets the bytecodes one at a time. Another kind of execution engine, one that is

faster but requires more memory, is a just-in-time compiler. In this scheme, the bytecodes of a method are compiled to native machine code the first time the method is invoked. The native machine code for the method is then cached, so it can be re-used the next time that same method is invoked. A third type of execution engine is an adaptive optimizer. In this approach, the virtual machine starts by interpreting bytecodes, but monitors the activity of the running program and identifies the most heavily used areas of code. As the program runs, the virtual machine compiles to native and optimizes just these heavily used areas. The rest of the code, which is not heavily used, remain as bytecodes which the virtual machine continues to interpret.

## 6.2 DVM and JIT

In Android OS lots of native code does the “heavy lifting” (e.g. graphics, media) out of the box. Also, JNI is available for developers to speed up their applications. On the other hand, while developing for Android OS it is unavoidable to use Java API. So, in some applications JIT can significantly increase execution speed. It is important to mention that JIT doesn’t improve speed if code already is doing work in native code.

One of the most important requirement for DVM is to use small amount of RAM. Therefore, light usage of memory is very important for DVM JIT design. In order to reduce memory usage, DVM use a trace-based JIT compilation approach – detects what are the hottest parts of code and compiles it into native code.

Since JIT at runtime increases system load and therefore battery usage, DVM JIT compiler does a lot of work upfront run time, in order to be fast at run time. Some optimizations are done on the development PC which has more power available. Also, some compilation are performed while device is charging.

## 7 Conclusion

The Dalvik VM has technical advantages over the Java VM for mobile environments, most notably aggressive use of copy-on-write memory sharing, so the entire VM and standard class library is shared among all Android SDK app processes, reducing the net per-process memory footprint.

## References

- [1] G.M. Amdahl, G.A. Blaauw, and FP Brooks. Architecture of the ibm system/360. *IBM Journal of Research and Development*, 8(2):87–101, 1964.
- [2] D. Bornstein. Dalvik vm internals. In *Google I/O Developer Conference*, volume 23, pages 17–30, 2008.



- [3] Dan Bornstein. Dalvik jit. <http://android-developers.blogspot.it/2010/05/dalvik-jit.html>, May 2010.
- [4] Jon Fingas. Idc: Android has a heady 59 percent of world smartphone share, iphone still on the way up. <http://www.engadget.com/2012/05/24/idc-q1-2012-world-smartphone-share/>, May 2012.
- [5] Google inc. Dalvik technical information. <http://source.android.com/tech/dalvik/index.html>, 2008.
- [6] Google Inc. Philosophy and goals, 2008.
- [7] Oracle Inc. About the java technology. <http://docs.oracle.com/javase/tutorial/getStarted/intro/definition.html>.
- [8] Oracle Inc. The class file format. <http://docs.oracle.com/javase/specs/jvms/se7/html/jvms-4.html>.
- [9] Oracle Inc. Java community process. <http://www.jcp.org/en/home/index>.
- [10] T. Lindholm and F. Yellin. *Java virtual machine specification*. Addison-Wesley Longman Publishing Co., Inc., 1999.
- [11] Frank Maker. Android dalvik virtual machine. <http://jaxenter.com/android-dalvik-virtual-machine-35498.html>, April 2011.
- [12] Se Hoon Park. Understanding jvm internals. <http://www.cubrid.org/blog/dev-platform/understanding-jvm-internals/>, May 2012.
- [13] The Android Open Source Project. Dalvik executable format. <http://source.android.com/tech/dalvik/dex-format.html>.
- [14] J.E. Smith and R. Nair. The architecture of virtual machines. *Computer*, 38(5):32–38, 2005.
- [15] Various. Java coffee. [http://en.wikipedia.org/wiki/Java\\_coffee](http://en.wikipedia.org/wiki/Java_coffee).
- [16] Various. Java (programming language). [http://en.wikipedia.org/wiki/Java\\_programming\\_language#History](http://en.wikipedia.org/wiki/Java_programming_language#History).
- [17] B. Venners. *Inside the Java virtual machine*. McGraw-Hill, Inc., 1996.
- [18] Wikipedia. List of jvm languages. [http://en.wikipedia.org/wiki/List\\_of\\_JVM\\_languages](http://en.wikipedia.org/wiki/List_of_JVM_languages).
- [19] Wikipedia. Virtual machine. [http://en.wikipedia.org/wiki/Virtual\\_machine](http://en.wikipedia.org/wiki/Virtual_machine).

## A Comparison of JVM and DVM

Table 2: Main features of JVM and DVM.

Feature	JVM	DVM	Note
JIT	+	+	In Dalvik JIT was introduced in Android 2.2
Bytecode	javac	Dalvik VM specifi	.class files recompiled with dx tool
Size of VM instructions	1 byte	2 Bytes	Increased expressiveness
Method references	No limit	65536	It is possible to load multiple dec files into environment, but is important only for really huge apps
where is code?	Multiple .class files	one .dex file	There is a trick to have few dec files.*
Memory consumption	Separate JVM, a lot of Ram	Shared libraries, small amount of ram	Dalvik is optimized for constraints
Security	In one VM every app can access everything	Separate VMs	very good